

# 强化学习基础及其在大语言模型中的应用

王榕甄

2025 年 7 月 3 日





# 提纲

- 强化学习基本概念及发展历程
- 策略梯度方法
- 策略改进方法
- 大语言模型中的强化对齐



## 强化学习的基本概念：三个随机变量、两个概率

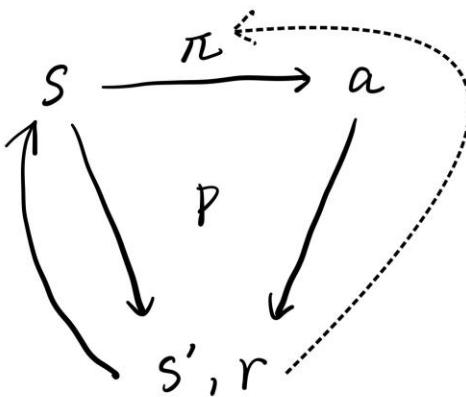
- “强化学习是一种从个体和环境交互的经验中学习如何将状态映射到动作，以获得最大奖励的学习机制。”
- 状态  $s$ , 动作  $a$ , 奖励  $r$
- 个体 (agent)：策略  $\pi(a|s)$
- 环境 (environment): 转移概率  $p(s', r|s, a)$
- 经过个体与环境  $t = 0, 1, \dots, T - 1$  次交互，形成轨迹 (trajectory or episode)

$$\tau = s_0, a_0, r_1, s_1, a_1, \dots, s_{T-1}, a_{T-1}, r_T, s_T$$

- 其中  $(s, a, r, s')$  为一步转移 (one transition step); 当  $t > T$  时,  $r_t = 0$

## 强化学习的问题建模

- 建模：马尔可夫决策过程 (Markov Decision Process, 或随机序贯决策, 或随机动态规划)



- 目标：找到一个好策略以最大化期望累积（折扣）奖励

$$J(\pi) = \mathbb{E}_{\tau \sim p_{\pi}(\tau)} \left[ \sum_{t=0}^{\infty} \gamma^t r_{t+1} \right], \quad \max_{\pi} J(\pi)$$

为什么要折扣？

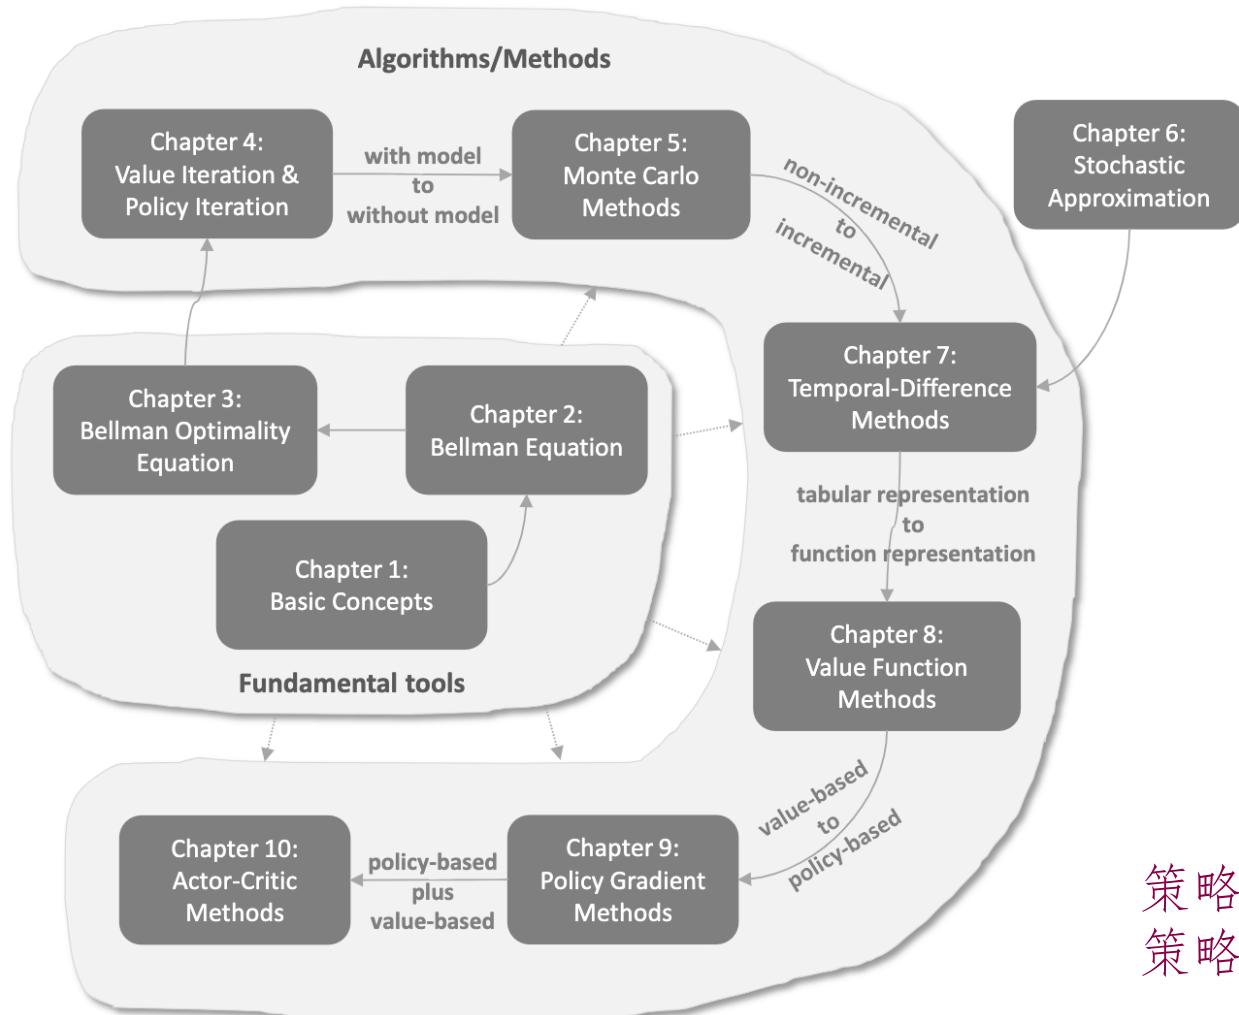
- $p_{\pi}(\tau) = p_{\pi}(s_0, a_0, r_1, s_1, \dots) = p(s_0) \prod_{t=0}^{\infty} \pi(a_t | s_t) p(s_{t+1}, r_{t+1} | s_t, a_t)$



## 强化学习的基本概念：价值函数

- 如何衡量某一状态下执行某一动作的好坏？
- 动作价值函数  $Q_{\pi}(s_t, a_t) = \mathbb{E}_{r_{t+1}, s_{t+1}, \dots \sim p_{\pi}(\cdot | s_t, a_t)} [\sum_{k=0}^{\infty} \gamma^k r_{t+k+1}]$
- 状态价值函数  $V_{\pi}(s_t) = \mathbb{E}_{a_t, r_{t+1}, s_{t+1}, \dots \sim p_{\pi}(\cdot | s_t)} [\sum_{k=0}^{\infty} \gamma^k r_{t+k+1}] = \mathbb{E}_{a_t \sim \pi(\cdot | s_t)} [Q_{\pi}(s_t, a_t)]$
- 优势函数  $A_{\pi}(s_t, a_t) = Q_{\pi}(s_t, a_t) - V_{\pi}(s_t)$

# 强化学习的发展历程



策略梯度方法  
策略改进方法



## 策略梯度方法 (Policy Gradient Methods)

- 回顾:  $\max_{\pi} J(\pi) = \mathbb{E}_{\tau \sim p_{\pi}(\tau)} [\sum_{t=0}^{\infty} \gamma^t r_{t+1}]$

- 参数化策略函数:  $\pi_{\theta}$

- 目标:

$$\max_{\theta} J(\theta) = \mathbb{E}_{\tau \sim p_{\pi_{\theta}}(\tau)} \left[ \sum_{t=0}^{\infty} \gamma^t r_{t+1} \right]$$

- 策略梯度方法:  $\theta \leftarrow \theta + \alpha \nabla_{\theta} J(\theta)$



# 策略梯度

$$\nabla_{\theta} J(\theta) = \mathbb{E}_{\tau \sim p_{\pi_{\theta}}(\tau)} \left[ \left( \sum_{t=0}^{\infty} \nabla_{\theta} \log \pi_{\theta}(a_t | s_t) \right) \left( \sum_{t=0}^{\infty} \gamma^t r_{t+1} \right) \right]$$

## Step 1: Log-Derivative Trick

- $J(\theta) = \mathbb{E}_{\tau \sim p_{\pi_{\theta}}(\tau)} [\sum_{t=0}^{\infty} \gamma^t r_{t+1}]$
- 令  $f(\tau) \triangleq \sum_{t=0}^{\infty} \gamma^t r_{t+1}$
- 则

## Step 2: Unrolling Log Probability

- $p_{\pi}(\tau) = p(s_0) \prod_{t=0}^{\infty} \pi(a_t | s_t) p(s_{t+1}, r_{t+1} | s_t, a_t)$
- 则



## REINFORCE 算法

$$\nabla_{\theta} J(\theta) = \mathbb{E}_{\tau \sim p_{\pi_{\theta}}(\tau)} \left[ \left( \sum_{t=0}^{\infty} \nabla_{\theta} \log \pi_{\theta}(a_t | s_t) \right) \left( \sum_{t=0}^{\infty} \gamma^t r_{t+1} \right) \right]$$

- REINFORCE algorithm:



1. sample  $\{\tau^i\}$  from  $\pi_{\theta}(\mathbf{a}_t | \mathbf{s}_t)$  (run the policy)
2.  $\nabla_{\theta} J(\theta) \approx \sum_i (\sum_t \nabla_{\theta} \log \pi_{\theta}(\mathbf{a}_t^i | \mathbf{s}_t^i)) (\sum_t r(\mathbf{s}_t^i, \mathbf{a}_t^i))$
3.  $\theta \leftarrow \theta + \alpha \nabla_{\theta} J(\theta)$

缺点：方差高、更新慢

- REINFORCE 的变体

$$V_{\pi}(s_t) = \mathbb{E}_{a_t, r_{t+1}, s_{t+1}, \dots \sim p_{\pi}(\cdot | s_t)} [\sum_{k=0}^{\infty} \gamma^k r_{t+k+1}]$$

- REINFORCE with Baseline:  $\sum_{t=0}^{\infty} \gamma^t r_{t+1} - V_{\pi}^{\phi}(s_0)$  无偏，优势函数
- Advantage Actor-Critic (A2C):  $r_1 + \gamma V_{\pi}^{\phi}(s_1) - V_{\pi}^{\phi}(s_0)$  很可能有偏
- Actor-Critic with GAE 结合以上二者



# 广义优势估计 (General Advantage Estimation, GAE)

- $\sum_{k=0}^{\infty} \gamma^{k+t} r_{t+k+1} - V_{\pi}(s_t)$
- $r_{t+1} + \gamma V_{\pi}(s_{t+1}) - V_{\pi}(s_t)$
- 观察:  $\hat{A}_t^{(1)} = -V(s_t) + r_t + \gamma V(s_{t+1}) = \delta_t^V$   
 $\hat{A}_t^{(2)} = -V(s_t) + r_t + \gamma r_{t+1} + \gamma^2 V(s_{t+2}) = \delta_t^V + \gamma \delta_{t+1}^V$   
 $\hat{A}_t^{(3)} = -V(s_t) + r_t + \gamma r_{t+1} + \gamma^2 r_{t+2} + \gamma^3 V(s_{t+3}) = \delta_t^V + \gamma \delta_{t+1}^V + \gamma^2 \delta_{t+2}^V$
- 加权平均:  $\hat{A}_t^{\text{GAE}(\gamma, \lambda)} := (1 - \lambda) \left( \hat{A}_t^{(1)} + \lambda \hat{A}_t^{(2)} + \lambda^2 \hat{A}_t^{(3)} + \dots \right)$   
 $= (1 - \lambda) (\delta_t^V + \lambda (\delta_t^V + \gamma \delta_{t+1}^V) + \lambda^2 (\delta_t^V + \gamma \delta_{t+1}^V + \gamma^2 \delta_{t+2}^V) + \dots)$   
 $= (1 - \lambda) (\delta_t^V (1 + \lambda + \lambda^2 + \dots) + \gamma \delta_{t+1}^V (\lambda + \lambda^2 + \lambda^3 + \dots)$   
 $\quad + \gamma^2 \delta_{t+2}^V (\lambda^2 + \lambda^3 + \lambda^4 + \dots) + \dots)$   
 $= (1 - \lambda) \left( \delta_t^V \left( \frac{1}{1 - \lambda} \right) + \gamma \delta_{t+1}^V \left( \frac{\lambda}{1 - \lambda} \right) + \gamma^2 \delta_{t+2}^V \left( \frac{\lambda^2}{1 - \lambda} \right) + \dots \right)$   
 $= \sum_{l=0}^{\infty} (\gamma \lambda)^l \delta_{t+l}^V$



## 广义优势估计 (General Advantage Estimation, GAE)

- GAE:  $\hat{A}_t^{\text{GAE}(\gamma, \lambda)} := \sum_{l=0}^{\infty} (\gamma \lambda)^l \delta_{t+l}^V$   
 $\text{GAE}(\gamma, 0) : \hat{A}_t := \delta_t = r_t + \gamma V(s_{t+1}) - V(s_t)$

$$\text{GAE}(\gamma, 1) : \hat{A}_t := \sum_{l=0}^{\infty} \gamma^l \delta_{t+l} = \sum_{l=0}^{\infty} \gamma^l r_{t+l} - V(s_t)$$

- 价值函数的估计与优化:

- 参数化:  $V_{\pi}(s_t) \approx V_{\pi}^{\phi}(s_t)$
- 优化:  $\min_{\phi} \frac{1}{N|\tau^i|} \sum_{i=1}^N \left\| V_{\pi}^{\phi}(s_t^i) - \hat{R}_t^i \right\|_2^2, \hat{R}_t = \sum_{k=0}^{T-t-1} \gamma^{k+t} r_{t+k+1}$

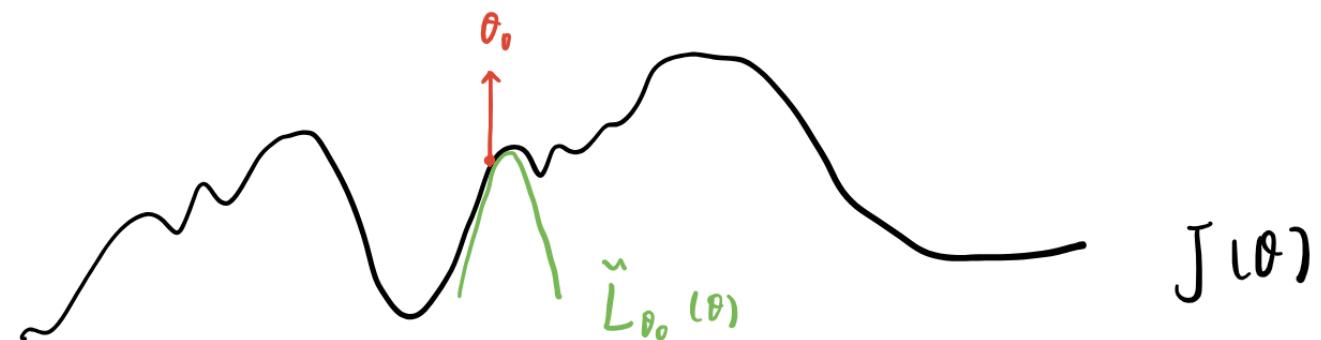
## 策略改进方法 (Policy Improvement Methods)

- 策略梯度方法：梯度上升，用一阶泰勒展开近似指导目标函数上升

$$\theta_1 \leftarrow \theta_0 + \alpha \nabla_{\theta} J(\theta)|_{\theta=\theta_0} \quad \Rightarrow \quad J(\theta_1) \geq J(\theta_0) ?$$

- 如何能保证目标函数上升？
- 考虑一个替代函数  $\tilde{L}(\theta)$ ，满足  $\tilde{L}(\theta) \leq J(\theta)$  且  $\tilde{L}(\theta_0) = J(\theta_0)$

$$\theta_1 \leftarrow \operatorname{argmax}_{\theta} \tilde{L}(\theta) \quad \Rightarrow \quad J(\theta_1) \geq \tilde{L}(\theta_1) \geq \tilde{L}(\theta_0) = J(\theta_0)$$





## 替代函数的构造

- 第一步：基于  $J(\theta_0)$  的  $J(\theta)$

$$J(\theta) = J(\theta_0) + \mathbb{E}_{s \sim \rho_{\pi_\theta}} \left[ \mathbb{E}_{a \sim \pi_\theta(\cdot|s)} [A_{\pi_{\theta_0}}(s, a)] \right]$$

- 其中  $\rho_\pi(s) = \Pr(s_0 = s | \pi, p) + \gamma \Pr(s_1 = s | \pi, p) + \gamma^2 \Pr(s_2 = s | \pi, p) + \dots$
- 推导：

- **Step 1 :**  $J(\theta) = J(\theta_0) + \mathbb{E}_{\tau \sim p_{\pi_\theta}(\tau)} \left[ \sum_{t=0}^{\infty} \gamma^t A_{\pi_{\theta_0}}(s_t, a_t) \right]$

- **Step 2 :**  $\mathbb{E}_{\tau \sim p_{\pi_\theta}(\tau)} \left[ \sum_{t=0}^{\infty} \gamma^t A_{\pi_{\theta_0}}(s_t, a_t) \right] = \mathbb{E}_{s \sim \rho_{\pi_\theta}} \left[ \mathbb{E}_{a \sim \pi_\theta(\cdot|s)} [A_{\pi_{\theta_0}}(s, a)] \right]$

$$\triangleright \max_{\pi} J(\pi) = \mathbb{E}_{\tau \sim p_{\pi}(\tau)} [\sum_{t=0}^{\infty} \gamma^t r_{t+1}]$$

# 替代函数的构造

Step 1.  $J(\theta) = J(\theta_0) + E_{\tau \sim \pi_\theta} \left[ \sum_{t=0}^{\infty} \gamma^t A_{\pi_{\theta_0}}(s_t, a_t) \right]$

证明:

$$\cdot A_{\pi_{\theta_0}}(s_t, a_t) = Q_{\pi_{\theta_0}}(s_t, a_t) - V_{\pi_{\theta_0}}(s_t)$$

$$= E_{r_{t+1}, s_{t+1}, \dots, v_{\pi_{\theta_0}}, |s_t, a_t|} \left[ \sum_{k=0}^{\infty} \gamma^k r_{t+k+1} \right] - V_{\pi_{\theta_0}}(s_t)$$

$$= E_{r_{t+1}, s_{t+1} \sim P(r_{t+1}, s_{t+1} | s_t, a_t)} \left[ r_{t+1} + \gamma V_{\pi_{\theta_0}}(s_{t+1}) \right] - V_{\pi_{\theta_0}}(s_t)$$

$$\cdot E_{\tau \sim \pi_\theta} \left[ \sum_{t=0}^{\infty} \gamma^t A_{\pi_{\theta_0}}(s_t, a_t) \right] = \sum_{t=0}^{\infty} \overbrace{E_{\tau \sim \pi_\theta} \left[ E_{r_{t+1}, s_{t+1} \sim P(r_{t+1}, s_{t+1} | s_t, a_t)} [r_{t+1} + \gamma V_{\pi_{\theta_0}}(s_{t+1})] - V_{\pi_{\theta_0}}(s_t) \right]}^{r^t}$$

$$= \sum_{t=0}^{\infty} \gamma^t E_{s_0, \dots, s_t, a_t, r_{t+1}, s_{t+1} \sim \pi_\theta} \left[ r_{t+1} + \gamma V_{\pi_{\theta_0}}(s_{t+1}) - V_{\pi_{\theta_0}}(s_t) \right]$$

$$= \sum_{t=0}^{\infty} \gamma^t E_{\tau \sim \pi_\theta} \left[ r_{t+1} + \gamma V_{\pi_{\theta_0}}(s_{t+1}) - V_{\pi_{\theta_0}}(s_t) \right]$$

$$= E_{\tau \sim \pi_\theta} \left[ r_1 + \gamma \cancel{V_{\pi_{\theta_0}}(s_1)} - V_{\pi_{\theta_0}}(s_0) \right. \\ \left. + \gamma r_2 + \gamma \cancel{V_{\pi_{\theta_0}}(s_2)} - \gamma \cancel{V_{\pi_{\theta_0}}(s_1)} \right. \\ \left. + \dots \right]$$

$$= E_{\tau \sim \pi_\theta} \left[ \sum_{t=0}^{\infty} \gamma^t r_{t+1} - \gamma V_{\pi_{\theta_0}}(s_0) \right]$$

$$= E_{\tau \sim \pi_\theta} \left[ \sum_{t=0}^{\infty} \gamma^t r_{t+1} \right] - E_{s_0 \sim P(s_0)} [V_{\pi_{\theta_0}}(s_0)]$$

$$= J(\theta) - J(\theta_0)$$

Step 2.  $E_{\tau \sim \pi_\theta} \left[ \sum_{t=0}^{\infty} \gamma^t A_{\pi_{\theta_0}}(s_t, a_t) \right] = \sup_{\pi \sim \pi_\theta} E_{\tau \sim \pi( \cdot | s )} [A_{\pi_{\theta_0}}(s, a)]$

证明:

$$\cdot E_{\tau \sim \pi_\theta} \left[ \sum_{t=0}^{\infty} \gamma^t A_{\pi_{\theta_0}}(s_t, a_t) \right] = \sum_{t=0}^{\infty} \gamma^t E_{\underbrace{s_0, \dots, s_t, a_t, \dots, v_{\pi_{\theta_0}}}_{t < t} [A_{\pi_{\theta_0}}(s_t, a_t)]}$$

$$= \sum_{t=0}^{\infty} \gamma^t \sum_{t < t} \sum_{s_t, a_t} P_{\pi_\theta}(s_t, a_t | s_t, a_t) [A_{\pi_{\theta_0}}(s_t, a_t)]$$

$$= \sum_{t=0}^{\infty} \gamma^t \sum_{s_t, a_t} P_{\pi_\theta}(s_t, a_t | s_t, a_t) [A_{\pi_{\theta_0}}(s_t, a_t)]$$

$$= \sum_{t=0}^{\infty} \gamma^t \sum_s \sum_a P_{\pi_\theta}(s_t = s, a_t = a) [A_{\pi_{\theta_0}}(s, a)] \quad \xrightarrow{\pi_\theta \text{ 和 } A_{\pi_{\theta_0}} \text{ 不区分时刻}}$$

$$= \sum_{t=0}^{\infty} \gamma^t \sum_s \sum_a P_{\pi_\theta}(s_t = s) \pi_\theta(s | a) [A_{\pi_{\theta_0}}(s, a)]$$

$$= \sum_s \underbrace{\left( \sum_{t=0}^{\infty} \gamma^t P_{\pi_\theta}(s_t = s) \right)}_{\triangleq P_{\pi_\theta}(s)} \sum_a \pi_\theta(s | a) [A_{\pi_{\theta_0}}(s, a)]$$

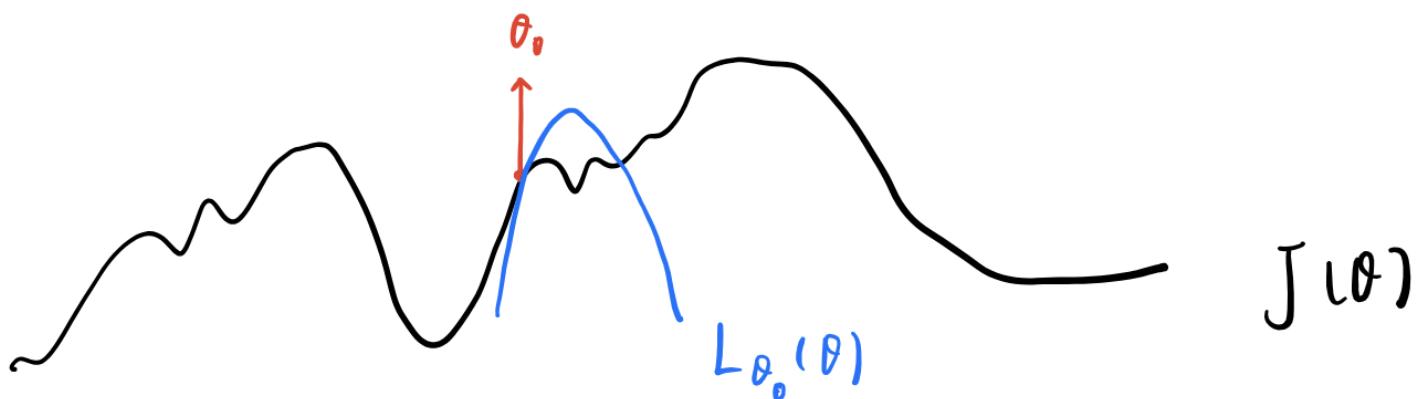
$$= \sum_s \sup_{\pi \sim \pi_\theta} E_{\tau \sim \pi( \cdot | s )} [A_{\pi_{\theta_0}}(s, a)]$$

# 替代函数的构造

- 第二步：近似

$$J(\theta) \approx J(\theta_0) + \mathbb{E}_{s \sim \rho_{\pi_{\theta_0}}} \left[ \mathbb{E}_{a \sim \pi_{\theta}(\cdot|s)} [A_{\pi_{\theta_0}}(s, a)] \right] = L(\theta)$$

- 满足条件： $L(\theta_0) = J(\theta_0)$  且  $\nabla_{\theta} L(\theta)|_{\theta=\theta_0} = \nabla_{\theta} J(\theta)|_{\theta=\theta_0}$



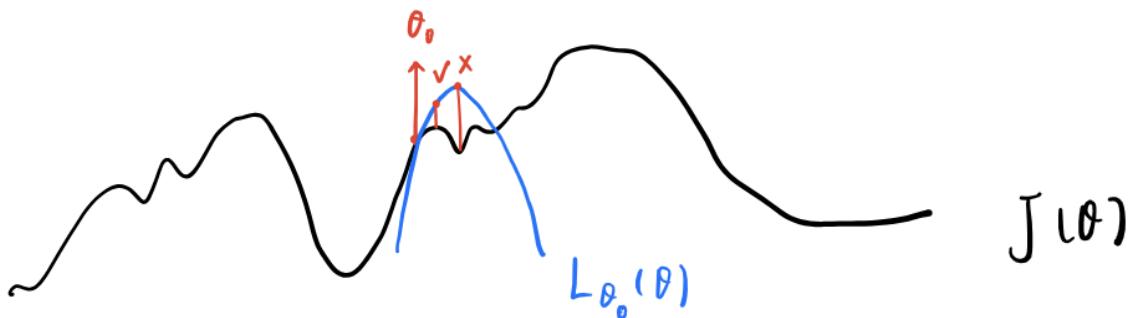
# 替代函数的构造

- 第三步：下界定理

$$J(\theta) \geq L(\theta) - C \max_s D_{\text{KL}} \left( \pi_{\theta_0}(\cdot|s) \middle\| \pi_{\theta}(\cdot|s) \right) = \tilde{L}(\theta)$$

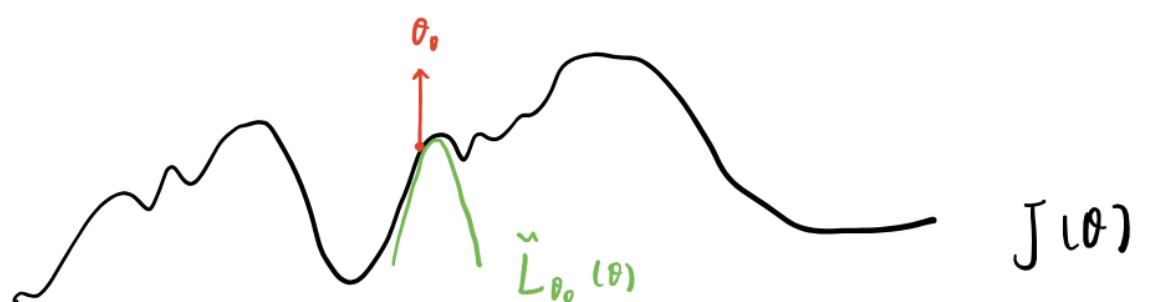
- $L(\theta) \triangleq J(\theta_0) + \mathbb{E}_{s \sim \rho_{\pi_{\theta}}} \left[ \mathbb{E}_{a \sim \pi_{\theta}(\cdot|s)} [A_{\pi_{\theta_0}}(s, a)] \right], \quad C \triangleq \frac{4\gamma \max_{s,a} |A_{\pi_{\theta_0}}(s, a)|}{(1-\gamma)^2}$
- $\rho_{\pi}(s) \triangleq \Pr(s_0 = s|\pi, p) + \gamma \Pr(s_1 = s|\pi, p) + \gamma^2 \Pr(s_2 = s|\pi, p) + \dots$

$$L(\theta_0) = J(\theta_0) \text{ 且 } \nabla_{\theta} L(\theta)|_{\theta=\theta_0} = \nabla_{\theta} J(\theta)|_{\theta=\theta_0}$$



$$J(\theta)$$

$$\tilde{L}(\theta) \leq J(\theta) \text{ 且 } \tilde{L}(\theta_0) = J(\theta_0)$$



$$J(\theta)$$



## 信赖域策略优化 (Trust Region Policy Optimization, TRPO)

- 有保证的目标:  $\max_{\theta} L(\theta) - C \max_s D_{\text{KL}} (\pi_{\theta_0}(\cdot|s) \| \pi_{\theta}(\cdot|s)) \gg$  步长过小!

- 替代地, 我们优化:

信赖域约束

$$\max_{\theta} L(\theta), \quad \text{subject to } \max_s D_{\text{KL}} (\pi_{\theta_0}(\cdot|s) \| \pi_{\theta}(\cdot|s)) \leq \delta$$

- 进一步替代地, 我们优化:

$$\max_{\theta} L(\theta), \quad \text{subject to } \mathbb{E}_{s \sim \rho_{\pi_{\theta_0}}} [D_{\text{KL}} (\pi_{\theta_0}(\cdot|s) \| \pi_{\theta}(\cdot|s))] \leq \delta$$

# 采样估计

- 目标函数:  $\max_{\theta} L(\theta) = \max_{\theta} J(\theta_0) + \mathbb{E}_{s \sim \rho_{\pi_{\theta_0}}} \left[ \mathbb{E}_{a \sim \pi_{\theta}(\cdot|s)} [A_{\pi_{\theta_0}}(s, a)] \right]$   
 $= \max_{\theta} \mathbb{E}_{s \sim \rho_{\pi_{\theta_0}}} \left[ \mathbb{E}_{a \sim \pi_{\theta}(\cdot|s)} [A_{\pi_{\theta_0}}(s, a)] \right]$   
 $= \max_{\theta} \mathbb{E}_{s \sim \rho_{\pi_{\theta_0}}} \left[ \mathbb{E}_{a \sim \pi_{\theta_0}(\cdot|s)} \left[ \frac{\pi_{\theta}(a|s)}{\pi_{\theta_0}(a|s)} A_{\pi_{\theta_0}}(s, a) \right] \right]$

重要性采样 (Importance Sampling) 技巧

$$\text{其中 } \rho_{\pi}(s) = \Pr(s_0 = s|\pi, p) + \gamma \Pr(s_1 = s|\pi, p) + \gamma^2 \Pr(s_2 = s|\pi, p) + \dots$$

$$= \max_{\theta} \sum_{t=0}^{\infty} \mathbb{E}_{s_t, a_t \sim \pi_{\theta_0}} \left[ \frac{\pi_{\theta}(a_t|s_t)}{\pi_{\theta_0}(a_t|s_t)} \gamma^t A_{\pi_{\theta_0}}(s_t, a_t) \right]$$

蒙特卡洛采样

$$\approx \max_{\theta} \frac{1}{N} \sum_{i=1}^N \sum_{t=0}^{|\tau^i|-1} \left[ \frac{\pi_{\theta}(a_t^i|s_t^i)}{\pi_{\theta_0}(a_t^i|s_t^i)} \gamma^t A_{\pi_{\theta_0}}(s_t^i, a_t^i) \right]$$

- 约束:  $\mathbb{E}_{s \sim \rho_{\pi_{\theta_0}}} [D_{\text{KL}}(\pi_{\theta_0}(\cdot|s) \| \pi_{\theta}(\cdot|s))] \approx \frac{1}{N} \sum_{i=1}^N \sum_{t=0}^{|\tau^i|-1} [\gamma^t \widehat{D}_{\text{KL}}(\pi_{\theta_0}(\cdot|s_t^i) \| \pi_{\theta}(\cdot|s_t^i))]$

- 优势:  $\hat{A}_t^{\text{GAE}(\gamma, \lambda)} = \sum_{l=0}^{\infty} (\gamma \lambda)^l \delta_{t+l}^V \quad \delta_t^V = r_t + \gamma V(s_{t+1}) - V(s_t)$

GAE



# 实际算法

---

**Algorithm 1** Trust Region Policy Optimization
 

---

- 1: Input: initial policy parameters  $\theta_0$ , initial value function parameters  $\phi_0$
- 2: Hyperparameters: KL-divergence limit  $\delta$ , backtracking coefficient  $\alpha$ , maximum number of backtracking steps  $K$
- 3: **for**  $k = 0, 1, 2, \dots$  **do**
- 4:   Collect set of trajectories  $\mathcal{D}_k = \{\tau_i\}$  by running policy  $\pi_k = \pi(\theta_k)$  in the environment.
- 5:   Compute rewards-to-go  $\hat{R}_t$ .
- 6:   Compute advantage estimates,  $\hat{A}_t$  (using any method of advantage estimation) based on the current value function  $V_{\phi_k}$ .

**1. 做近似**——构造函数  $\tilde{L}$  近似目标函数  $J(\theta)$ :

- (a). 设当前策略网络参数是  $\theta_{\text{now}}$ 。用策略网络  $\pi(a | s; \theta_{\text{now}})$  控制智能体与环境交互，玩完一局游戏，记录下轨迹：

$$s_1, a_1, r_1, s_2, a_2, r_2, \dots, s_n, a_n, r_n.$$

- (b). 对于所有的  $t = 1, \dots, n$ ，计算折扣回报  $u_t = \sum_{k=t}^n \gamma^{k-t} \cdot r_k$ 。

- (c). 得出近似函数：

$$\tilde{L}(\theta | \theta_{\text{now}}) = \sum_{t=1}^n \frac{\pi(a_t | s_t; \theta)}{\pi(a_t | s_t; \theta_{\text{now}})} \cdot u_t.$$

**2. 最大化**——用某种数值算法求解带约束的最大化问题：

$$\theta_{\text{new}} = \operatorname{argmax}_{\theta} \tilde{L}(\theta | \theta_{\text{now}}); \quad \text{s.t. } \|\theta - \theta_{\text{now}}\|_2 \leq \Delta.$$

此处的约束条件是二范数距离。可以把它替换成 KL 散度，即公式 (9.10)。

$$\begin{aligned} & \max_{\theta} L(\theta) \\ &= \max_{\theta} \sum_{t=0}^{\infty} \mathbb{E}_{s_t, a_t \sim \pi_{\theta_0}} \left[ \frac{\pi_{\theta}(a_t | s_t)}{\pi_{\theta_0}(a_t | s_t)} A_{\pi_{\theta_0}}(s_t, a_t) \right] \\ & \approx \max_{\theta} \frac{1}{N} \sum_{i=1}^N \sum_{t=0}^{|\tau^i|-1} \left[ \frac{\pi_{\theta}(a_t^i | s_t^i)}{\pi_{\theta_0}(a_t^i | s_t^i)} \gamma^t A_{\pi_{\theta_0}}(s_t^i, a_t^i) \right] \end{aligned}$$

**Algorithm 1** Trust Region Policy Optimization

$$\max_{\pi} J(\pi) = \mathbb{E}_{\tau \sim p_{\pi}(\tau)} [\sum_{t=0}^{\infty} \gamma^t r_{t+1}]$$

# 实际算法

- 1: Input: initial policy parameters  $\theta_0$ , initial value function parameters  $\phi_0$
- 2: Hyperparameters: KL-divergence limit  $\delta$ , backtracking coefficient  $\alpha$ , maximum number of backtracking steps  $K$
- 3: **for**  $k = 0, 1, 2, \dots$  **do**
- 4:   Collect set of trajectories  $\mathcal{D}_k = \{\tau_i\}$  by running policy  $\pi_k = \pi(\theta_k)$  in the environment.
- 5:   Compute rewards-to-go  $\hat{R}_t$ .
- 6:   Compute advantage estimates,  $\hat{A}_t$  (using any method of advantage estimation) based on the current value function  $V_{\phi_k}$ .
- 7:   Estimate policy gradient as

$$\hat{g}_k = \frac{1}{|\mathcal{D}_k|} \sum_{\tau \in \mathcal{D}_k} \sum_{t=0}^T \nabla_{\theta} \log \pi_{\theta}(a_t | s_t)|_{\theta_k} \hat{A}_t.$$

- 8:   Use the conjugate gradient algorithm to compute

$$\hat{x}_k \approx \hat{H}_k^{-1} \hat{g}_k,$$

where  $\hat{H}_k$  is the Hessian of the sample average KL-divergence.

- 9:   Update the policy by backtracking line search with

$$\theta_{k+1} = \theta_k + \alpha^j \sqrt{\frac{2\delta}{\hat{x}_k^T \hat{H}_k \hat{x}_k}} \hat{x}_k,$$

where  $j \in \{0, 1, 2, \dots K\}$  is the smallest value which improves the sample loss and satisfies the sample KL-divergence constraint.

- 10:   Fit value function by regression on mean-squared error:

$$\phi_{k+1} = \arg \min_{\phi} \frac{1}{|\mathcal{D}_k| T} \sum_{\tau \in \mathcal{D}_k} \sum_{t=0}^T \left( V_{\phi}(s_t) - \hat{R}_t \right)^2,$$

typically via some gradient descent algorithm.

- 11: **end for**

$$\begin{aligned} & \max_{\theta} L(\theta) \\ &= \max_{\theta} \sum_{t=0}^{\infty} \mathbb{E}_{s_t, a_t \sim \pi_{\theta_0}} \left[ \frac{\pi_{\theta}(a_t | s_t)}{\pi_{\theta_0}(a_t | s_t)} A_{\pi_{\theta_0}}(s_t, a_t) \right] \\ &\approx \max_{\theta} \frac{1}{N} \sum_{i=1}^N \sum_{t=0}^{|\tau^i|-1} \left[ \frac{\pi_{\theta}(a_t^i | s_t^i)}{\pi_{\theta_0}(a_t^i | s_t^i)} \gamma^t A_{\pi_{\theta_0}}(s_t^i, a_t^i) \right] \end{aligned}$$



# 近端策略优化 (Proximal Policy Optimization, PPO)

- 回顾:

- 下界定理:  $J(\theta) \geq L(\theta) - CD_{\text{KL}}^{\max}(\pi_{\theta_0}, \pi_{\theta})$ ,  $C = \frac{4\gamma \max_{s,a} |A_{\pi_{\theta_0}}(s, a)|}{(1-\gamma)^2}$
- $\max_{\theta} L(\theta) = \max_{\theta} \sum_{t=0}^{\infty} \mathbb{E}_{s_t, a_t \sim \pi_{\theta_0}} \left[ \frac{\pi_{\theta}(a_t|s_t)}{\pi_{\theta_0}(a_t|s_t)} A_{\pi_{\theta_0}}(s_t, a_t) \right]$  s.t.  $\max_s D_{\text{KL}}(\pi_{\theta_0}(\cdot|s) \| \pi_{\theta}(\cdot|s)) \leq \delta$
- PPO-Penalty:  $\max_{\theta} L(\theta) - \beta D_{\text{KL}}^{\max}(\pi_{\theta_0}, \pi_{\theta})$
- PPO-Clip:

$$\max_{\theta} \sum_{t=0}^{\infty} \mathbb{E}_{s_t, a_t \sim \pi_{\theta_0}} \left[ \min \left( \frac{\pi_{\theta}(a_t|s_t)}{\pi_{\theta_0}(a_t|s_t)} A_{\pi_{\theta_0}}(s_t, a_t), \text{clip} \left( \frac{\pi_{\theta}(a_t|s_t)}{\pi_{\theta_0}(a_t|s_t)}, 1 - \epsilon, 1 + \epsilon \right) A_{\pi_{\theta_0}}(s_t, a_t) \right) \right]$$



## PPO-Clip

**优势为正：**假设该状态-动作对的优势为正，在这种情况下，其对目标的贡献减少为

$$L(s, a, \theta_k, \theta) = \min \left( \frac{\pi_{\theta}(a|s)}{\pi_{\theta_k}(a|s)}, (1 + \epsilon) \right) A^{\pi_{\theta_k}}(s, a).$$

由于优势为正，如果行动的可能性增大（即 $\pi_{\theta}(a|s)$ 增加），目标也会随之增加。但该项中的最小值限制了目标的增幅。一旦，最小值就会生效，该项就会达到的上限。因此：远离旧策略并不会给新策略带来好处。 $\pi_{\theta}(a|s) > (1 + \epsilon)\pi_{\theta_k}(a|s)(1 + \epsilon)A^{\pi_{\theta_k}}(s, a)$

**优势为负：**假设该状态-动作对的优势为负，在这种情况下，其对目标的贡献减少为

$$L(s, a, \theta_k, \theta) = \max \left( \frac{\pi_{\theta}(a|s)}{\pi_{\theta_k}(a|s)}, (1 - \epsilon) \right) A^{\pi_{\theta_k}}(s, a).$$

由于优势为负，如果行动发生的可能性降低（即 $\pi_{\theta}(a|s)$ 减小），目标就会增加。但该项中的最大值限制了目标的增幅。一旦，最大值就会生效，该项就会达到的上限。因此，再次强调：远离旧策略并不会给新策略带来好处。 $\pi_{\theta}(a|s) < (1 - \epsilon)\pi_{\theta_k}(a|s)(1 - \epsilon)A^{\pi_{\theta_k}}(s, a)$

到目前为止，我们所看到的是，裁剪通过消除政策发生剧烈变化的动机，起到了正则化的作用，而超参数 $\epsilon$ 对应于新政策与旧政策相差多远，同时仍然有利于实现目标。



# 实际算法

---

## Algorithm 1 PPO-Clip

---

- 1: Input: initial policy parameters  $\theta_0$ , initial value function parameters  $\phi_0$
- 2: **for**  $k = 0, 1, 2, \dots$  **do**
- 3:   Collect set of trajectories  $\mathcal{D}_k = \{\tau_i\}$  by running policy  $\pi_k = \pi(\theta_k)$  in the environment.
- 4:   Compute rewards-to-go  $\hat{R}_t$ .
- 5:   Compute advantage estimates,  $\hat{A}_t$  (using any method of advantage estimation) based on the current value function  $V_{\phi_k}$ .
- 6:   Update the policy by maximizing the PPO-Clip objective:

$$\theta_{k+1} = \arg \max_{\theta} \frac{1}{|\mathcal{D}_k|T} \sum_{\tau \in \mathcal{D}_k} \sum_{t=0}^T \min \left( \frac{\pi_{\theta}(a_t | s_t)}{\pi_{\theta_k}(a_t | s_t)} A^{\pi_{\theta_k}}(s_t, a_t), g(\epsilon, A^{\pi_{\theta_k}}(s_t, a_t)) \right),$$

typically via stochastic gradient ascent with Adam.

- 7:   Fit value function by regression on mean-squared error:

$$\phi_{k+1} = \arg \min_{\phi} \frac{1}{|\mathcal{D}_k|T} \sum_{\tau \in \mathcal{D}_k} \sum_{t=0}^T \left( V_{\phi}(s_t) - \hat{R}_t \right)^2,$$

typically via some gradient descent algorithm.

- 8: **end for**
-

# 大语言模型中的强化学习



# 大语言模型中的强化学习：从 RLHF 说起

- 问题建模：单步策略，学习 question 对应的 output

$$J(\theta) = \mathbb{E}_{\tau \sim p_{\pi_{\theta}}(\tau)} \left[ \sum_{t=0}^1 \gamma^t r_{t+1} \right] = \mathbb{E}_{s_0 \sim p(s_0), a_0 \sim \pi_{\theta}(\cdot | s_0)} [r_1] = \mathbb{E}_{q \sim p(q), o \sim \pi_{\theta}(\cdot | q)} [r_1]$$

- 奖励模型：

- Rule-based (or verifiable, RLVR)
- Model-based: 结果奖励模型 (outcome RM)、过程奖励模型 (process RM)

➤ 结果奖励模型  $r_{\psi}(q, o)$ , BT 公式 + 交叉熵损失

$$\mathcal{L}_R(r_{\phi}, \mathcal{D}) = -\mathbb{E}_{(x, y_w, y_l) \sim \mathcal{D}} [\log \sigma(r_{\phi}(x, y_w) - r_{\phi}(x, y_l))] \quad \mathcal{D} = \{x^{(i)}, y_w^{(i)}, y_l^{(i)}\}_{i=1}^N$$

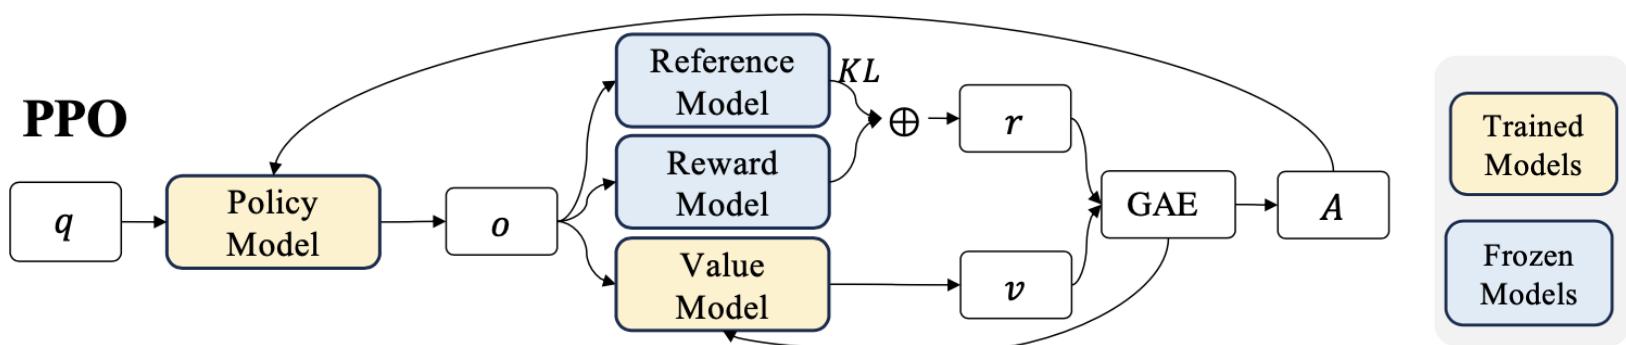
$$\text{➤ } r_1 = r(q, o) = r_{\psi}(q, o) - \beta \log \frac{\pi_{\theta}(\cdot | q)}{\pi_{\text{ref}}(\cdot | q)}$$

# 大语言模型中的强化学习：从 RLHF 说起

- 求解算法：PPO

$$J_{\text{PPO}}(\theta) = \mathbb{E}_{q \sim p(q), o \sim \pi_{\theta_0}(\cdot|q)} \left[ \min \left( \frac{\pi_{\theta}(o|q)}{\pi_{\theta_0}(o|q)} A_{\pi_{\theta_0}}^{\psi, \phi}(q, o), \text{clip} \left( \frac{\pi_{\theta}(o|q)}{\pi_{\theta_0}(o|q)}, 1 - \epsilon, 1 + \epsilon \right) A_{\pi_{\theta_0}}^{\psi, \phi}(q, o) \right) \right]$$

$$\hat{A}_t^{\text{GAE}(\gamma, \lambda)} := \sum_{l=0}^{\infty} (\gamma \lambda)^l \delta_{t+l}^V \quad \delta_t^V = r_t + \gamma V(s_{t+1}) - V(s_t)$$



奖励模型  $r_{\psi}(q, o)$  6B

价值模型  $V_{\pi}(s_t)$  6B

策略模型  $\pi_{\theta}(o|q)$  175B



## Group Relative Policy Optimization (GRPO)

- 价值模型占用内存、耗费计算量  $>>$  不使用价值模型估计优势
- 对给定  $q$ , 从  $\pi_{\theta_0}$  中采样一组回复  $\{o^1, o^2, \dots, o^G\}$ , 打分  $r = \{r^1, r^2, \dots, r^G\}$

$$\tilde{r}^i = \frac{r^i - \text{mean}(r)}{\text{std}(r)}, \quad A_{\pi_{\theta_0}}^{\psi}(q, o^i) = \tilde{r}^i$$

- $J_{\text{GRPO}}(\theta) = \mathbb{E}_{q \sim p(q), \{o^1, o^2, \dots, o^G\} \sim \pi_{\theta_0}(\cdot | q)}$

$$\left[ \frac{1}{G} \sum_{i=1}^G \min \left( \frac{\pi_{\theta}(o^i | q)}{\pi_{\theta_0}(o^i | q)} A_{\pi_{\theta_0}}^{\psi}(q, o^i), \text{clip} \left( \frac{\pi_{\theta}(o^i | q)}{\pi_{\theta_0}(o^i | q)}, 1 - \epsilon, 1 + \epsilon \right) A_{\pi_{\theta_0}}^{\psi}(q, o^i) \right) - \beta D_{\text{KL}}(\pi_{\theta}(\cdot | q) \| \pi_{\text{ref}}(\cdot | q)) \right]$$

# 拓展到多步策略

- PPO

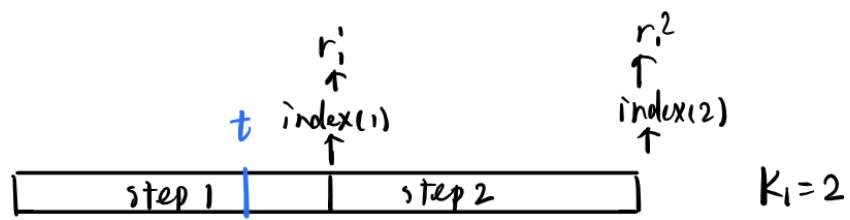
$$\mathcal{J}_{PPO}(\theta) = \mathbb{E}[q \sim P(Q), o \sim \pi_{\theta_{old}}(O|q)] \frac{1}{|o|} \sum_{t=1}^{|o|} \min \left[ \frac{\pi_{\theta}(o_t|q, o_{<t})}{\pi_{\theta_{old}}(o_t|q, o_{<t})} A_t, \text{clip} \left( \frac{\pi_{\theta}(o_t|q, o_{<t})}{\pi_{\theta_{old}}(o_t|q, o_{<t})}, 1 - \varepsilon, 1 + \varepsilon \right) A_t \right].$$

- GRPO

$$\mathcal{J}_{GRPO}(\theta) = \mathbb{E}[q \sim P(Q), \{o_i\}_{i=1}^G \sim \pi_{\theta_{old}}(O|q)]$$

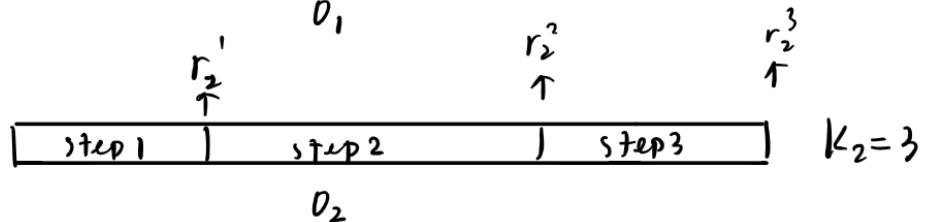
$$\frac{1}{G} \sum_{i=1}^G \frac{1}{|o_i|} \sum_{t=1}^{|o_i|} \left\{ \min \left[ \frac{\pi_{\theta}(o_{i,t}|q, o_{i,<t})}{\pi_{\theta_{old}}(o_{i,t}|q, o_{i,<t})} \hat{A}_{i,t}, \text{clip} \left( \frac{\pi_{\theta}(o_{i,t}|q, o_{i,<t})}{\pi_{\theta_{old}}(o_{i,t}|q, o_{i,<t})}, 1 - \varepsilon, 1 + \varepsilon \right) \hat{A}_{i,t} \right] - \beta \mathbb{D}_{KL} [\pi_{\theta} || \pi_{ref}] \right\}$$

- 过程奖励模型

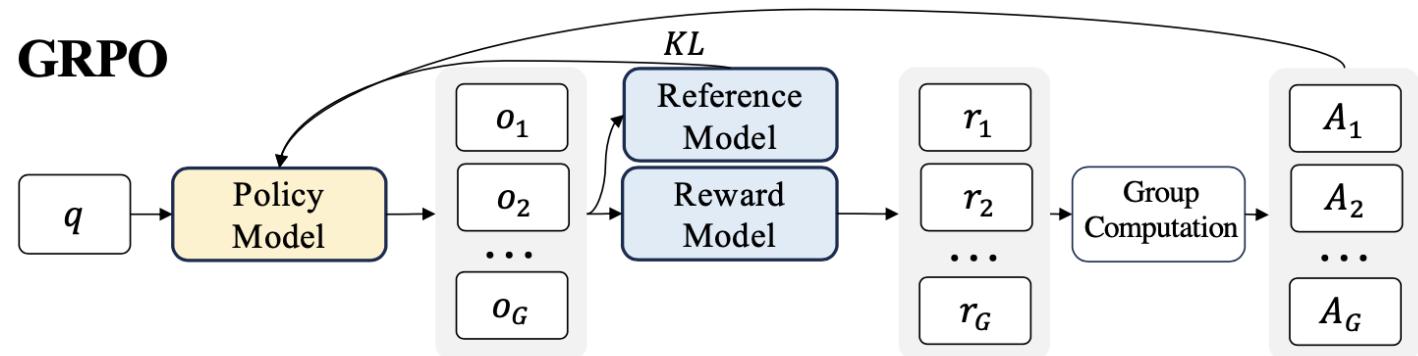


$$\tilde{r}_1^{\text{index}(1)} = \frac{r_1^1 - \text{mean}(r_1^1, r_1^2, r_2^1, r_2^2, r_2^3)}{\text{std}(r \dots)}$$

$$\hat{A}_{(q, o_1, \leq t)}^{\rho} = \tilde{r}_1^{\text{index}(1)} + \tilde{r}_1^{\text{index}(2)}$$



# 实际算法




---

## Algorithm 1 Iterative Group Relative Policy Optimization

---

**Input** initial policy model  $\pi_{\theta_{\text{init}}}$ ; reward models  $r_\varphi$ ; task prompts  $\mathcal{D}$ ; hyperparameters  $\varepsilon, \beta, \mu$

- 1: policy model  $\pi_\theta \leftarrow \pi_{\theta_{\text{init}}}$
- 2: **for** iteration = 1, ..., I **do**
- 3:   reference model  $\pi_{ref} \leftarrow \pi_\theta$
- 4:   **for** step = 1, ..., M **do**
- 5:     Sample a batch  $\mathcal{D}_b$  from  $\mathcal{D}$
- 6:     Update the old policy model  $\pi_{\theta_{old}} \leftarrow \pi_\theta$
- 7:     Sample  $G$  outputs  $\{o_i\}_{i=1}^G \sim \pi_{\theta_{old}}(\cdot | q)$  for each question  $q \in \mathcal{D}_b$
- 8:     Compute rewards  $\{r_i\}_{i=1}^G$  for each sampled output  $o_i$  by running  $r_\varphi$
- 9:     Compute  $\hat{A}_{i,t}$  for the  $t$ -th token of  $o_i$  through group relative advantage estimation.
- 10:    **for** GRPO iteration = 1, ...,  $\mu$  **do**
- 11:      Update the policy model  $\pi_\theta$  by maximizing the GRPO objective (Equation 21)
- 12:    Update  $r_\varphi$  through continuous training using a replay mechanism.

**Output**  $\pi_\theta$

---



# 总结

- 强化学习系统性很强，一环扣一环
- 目前大语言模型中的强化学习：

本质吗？是最好的吗？——算法还很粗糙、应用还很初步

- 学习资料：
  - 教材：
    - Richard S. Sutton and Andrew G. Barto, 2020, Reinforcement Learning: An Introduction
    - 王树森 张志华, 2022, 深度强化学习 (初稿)
  - 课程：
    - 西湖大学赵世钰【【强化学习的数学原理】课程：从零开始到透彻理解（完结）】
    - OpenAI, Welcome to Spinning Up in Deep RL!



# 总结

- 其它资料：
  - 博客： [Seita's Place: Going Deeper Into Reinforcement Learning: Fundamentals of Policy Gradients](#)
  - 博客： [Seita's Place: Notes on the Generalized Advantage Estimation Paper](#)
  - 博客： [Weng Lilian, A \(Long\) Peek into Reinforcement Learning, 2018](#)
  - 博客： [Weng Lilian, Policy Gradient Algorithms, 2018](#)
  - 课程： [Sergey Levine, CS 285 UC Berkeley, Deep Reinforcement Learning, 2023](#)
  - 教材： [Kevin Murphy, Reinforcement Learning: An Overview](#)
  - 课程： [David Silver, UCL, Reinforcement Learning, 2015](#)

**Thank you for listening!**